



Interpretable Machine Learning

What a ML project looks like?



Machine learning
students at the
beginning of a project

vs.

Machine learning
students at the end of
a project





Ali Izadi

Miras-tech - 1.5 yr


Ceit - 94



My experiences and failures in machine learning projects

Machine Learning Interpretation

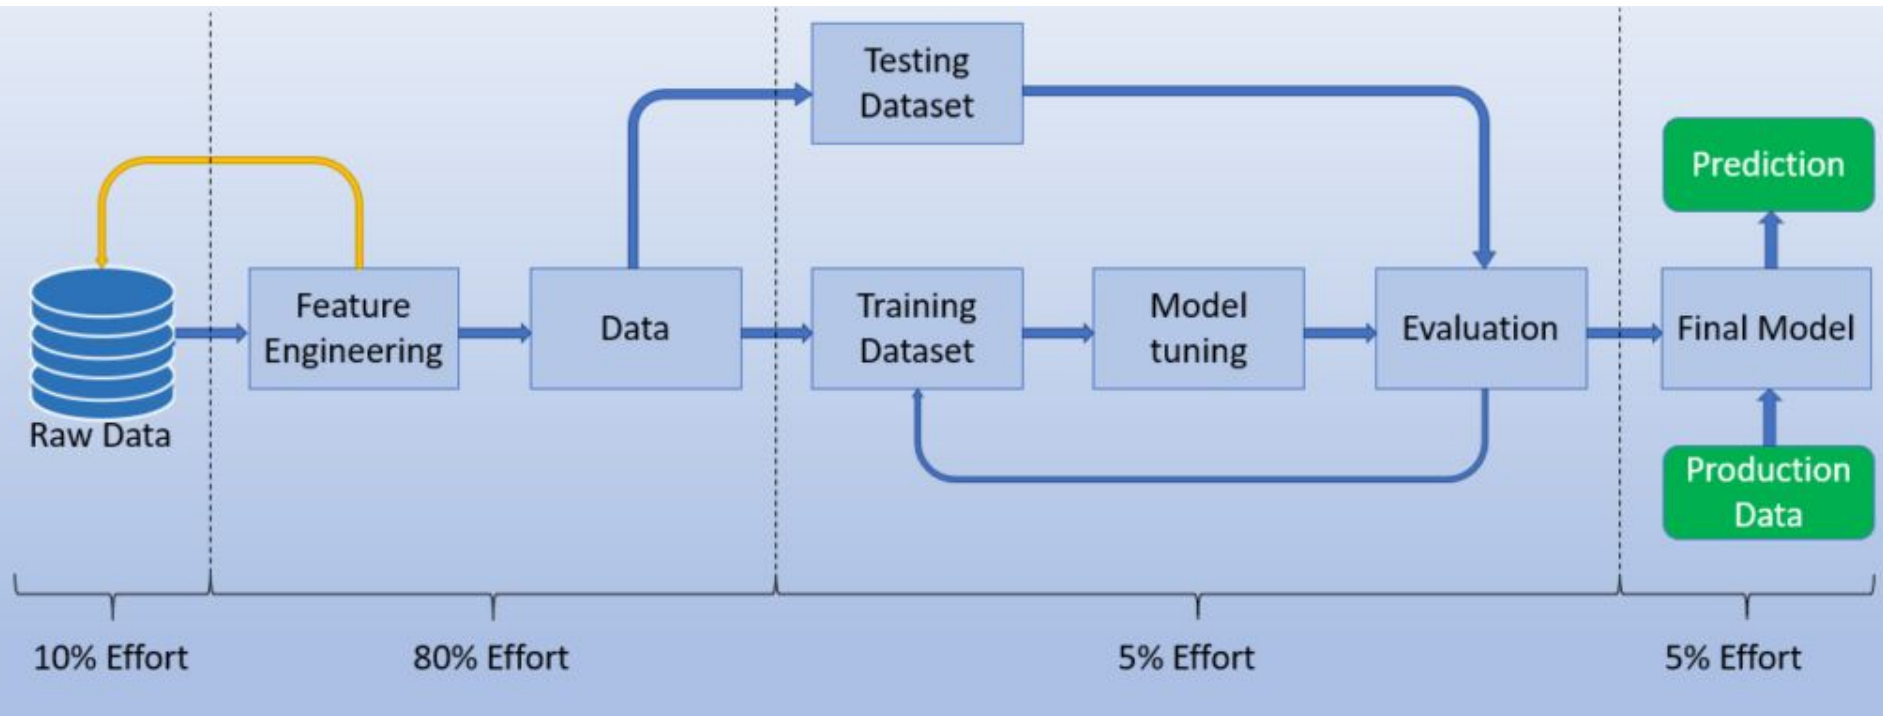
Field of Data Science in Iran



The importance of machine learning interpretation

Machine learning interpretation methods

Machine Learning Projects Steps:





Expectation vs Reality

Just learn machine learning models
and how to use tools such as tensorflow
or scikit-learn



Expectation vs Reality

Just accuracy or more advanced metrics
are important.



Expectation vs Reality

Data Science is just experiment without
innovation.



Expectation vs Reality

Most of us believe preparing data is the most time consuming part of data science but I believe feature engineering is the most important one.



Expectation vs Reality

You must learn more about the domain of
your problem than software engineers



Expectation vs Reality

PhD vs Msc vs Bsc



Interpretable Machine Learning

How can I trust your model?

How does your model really make its decision?





The importance of ML interpretation:

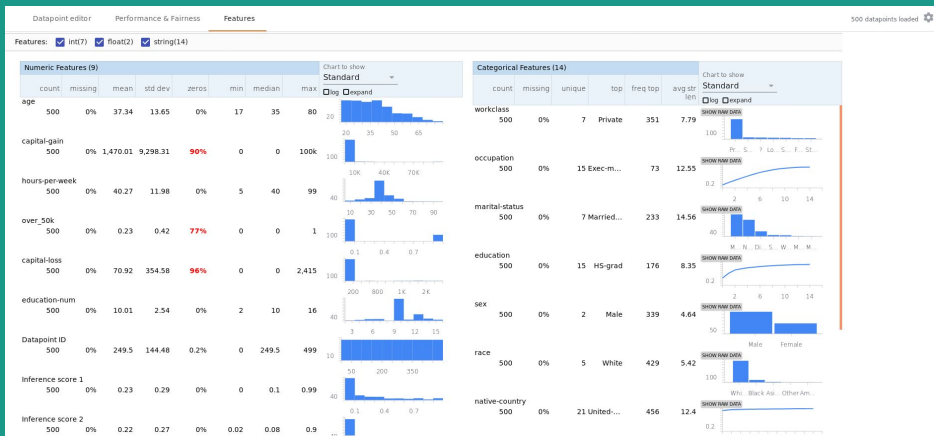
1 Reliability

3 Regulations

2 Simplicity

4 fairness

Google



What If...

you could inspect a machine learning model,
with minimal coding required?



People + AI Guidebook

Designing human-centered AI products

User Needs + Defining Success

Identify user needs, find AI opportunities,
and design your reward function.

Data Collection + Evaluation

Decide what data are required to meet your
user needs, source data, and tune your AI.

Mental Models

Introduce users to the AI system and set
expectations for system-change over time.

Explainability + Trust

Explain the AI system and determine if,
when, and how to show model confidence.

Feedback + Control

Design feedback and control mechanisms
to improve your AI and the user experience.

Errors + Graceful Failure

Identify and diagnose AI and context errors
and communicate the way forward.

McKinsey

Keys to building a data driven strategy:


Companies should repeatedly ask, “What’s the least complex model that would improve our performance?”

Leading your organization to responsible AI:

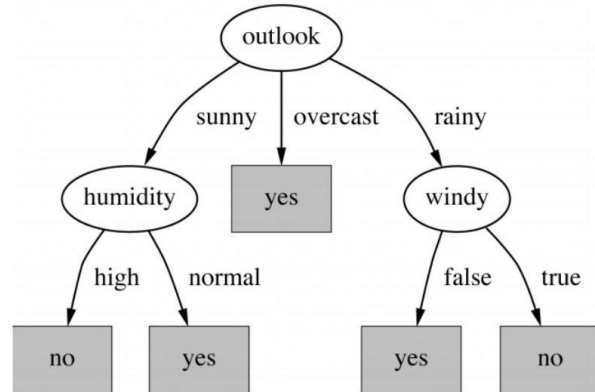
The ability to explain model outputs to stakeholders is a major lever in ensuring compliance with expanding regulatory and public expectations and in fostering trust to accelerate adoption.




Taxonomy of interpretability methods:



Model specific



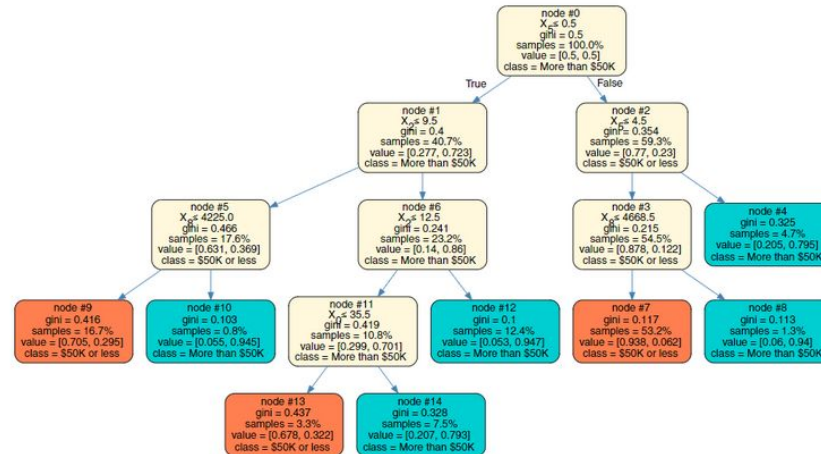
Decision Tree



Regression

Model agnostic:

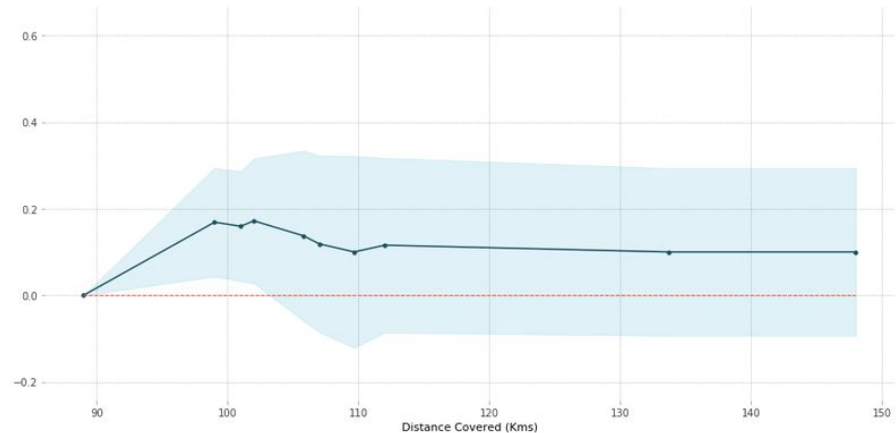
is there a way to build interpretable approximations of really complex models?



Permutation Importance:


1	165349.2	136897.8	471784.1	192261.83	0
2	162597.7	151377.59	443898.53	191792.06	1
3	153441.51	101145.55	407934.54	191050.39	1
...
48	0	135426.92	0	42559.73	1
49	542.05	51743.15	0	35673.41	0
50	0	116983.8	45173.06	14681.4	1

Partial Dependence Plot:




SHAP:


Why model predicts 0.7 instead of 0.5?



Interpretability for deep learning:



Saliency Map



Class Activation Map



Will Koehrsen

@koehrsen_will



Develop new algorithms as a PhD student: \$30k/year

Use pre-built sklearn models as a data scientist: \$120k/year

Build regression models in excel as a hedge fund analyst: \$200k/year

Make pie charts as a CEO: \$14 million/year

3:48 · 28 Oct 19 · [Twitter Web App](#)

5,969 Retweets **27.6K** Likes



Will Koehrsen @koehrsen_will · 28 Oct



On a more serious note, a complicated algorithm does not equal a better algorithm if it's not understood and trusted by people. Helping people make better decisions often means keeping models as simple as possible.



22



211



1,915





Data Science In Iran

A photograph of Donald Trump speaking at a podium. He is wearing a dark blue suit, a white shirt, and a red and blue striped tie. He is pointing his right index finger towards the left side of the frame. A microphone is visible in front of him. The background is dark and out of focus.

**YOU A DATA SCIENTIST,
YOU A DATA SCIENTIST**

EVERYONE'S A DATA SCIENTIST

Data Scientists



What my mom thinks I do



What my friends think I do



What society thinks I do



What investors think I do




What I think I do



What I really do

Future AI demand trajectory, % change in AI spending over next 3 years¹





Recommendations


- www.mckinsey.com
- www.pair.withgoogle.com
- Book: thinking with data, Max Shron
- Book: interpretable machine learning, Christoph Molnar
- Course: machine learning Explainability, [kaggle.com](https://www.kaggle.com)



Thank you.



aliizadi2030@gmail.com



@Aliizadi76